

```
1 // headers
2 #include <stdio.h>
3
4 #include <cuda.h> // for CUDA
5
6 // global variables
7 int inputLength=5;
8
9 float *hostInput1=NULL;
10 float *hostInput2=NULL;
11 float *hostOutput=NULL;
12
13 float *deviceInput1=NULL;
14 float *deviceInput2=NULL;
15 float *deviceOutput=NULL;
16
17 // global kernel function definition
18 __global__ void vecAdd(float *in1,float *in2,float *out,int len)
19 {
20     // variable declarations
21     int i=blockIdx.x * blockDim.x + threadIdx.x;
22     // code
23     if(i < len)
24     {
25         out[i]=in1[i]+in2[i];
26     }
27 }
28
29 int main(int argc,char *argv[])
30 {
31     // function declarations
32     void cleanup(void);
33
34     // code
35     // allocate host-memory
36     hostInput1=(float *)malloc(inputLength * sizeof(float));
37     if(hostInput1== NULL)
38     {
39         printf("CPU Memory Fatal Error = Can Not Allocate Enough Memory For
40             Host Input Array 1.\nExiting ...\n");
41         cleanup();
42         exit(EXIT_FAILURE);
43     }
44     hostInput2=(float *)malloc(inputLength * sizeof(float));
45     if(hostInput2== NULL)
46     {
47         printf("CPU Memory Fatal Error = Can Not Allocate Enough Memory For
48             Host Input Array 2.\nExiting ...\n");
49         cleanup();
50         exit(EXIT_FAILURE);
51     }
52     hostOutput=(float *)malloc(inputLength * sizeof(float));
53     if(hostOutput== NULL)
54     {
```

```
55     printf("CPU Memory Fatal Error = Can Not Allocate Enough Memory For
        Host Output Array.\nExiting ... \n");
56     cleanup();
57     exit(EXIT_FAILURE);
58 }
59
60 // fill above input host vectors with arbitrary but hard-coded data
61 hostInput1[0]=101.0;
62 hostInput1[1]=102.0;
63 hostInput1[2]=103.0;
64 hostInput1[3]=104.0;
65 hostInput1[4]=105.0;
66
67 hostInput2[0]=201.0;
68 hostInput2[1]=202.0;
69 hostInput2[2]=203.0;
70 hostInput2[3]=204.0;
71 hostInput2[4]=205.0;
72
73 // allocate device-memory
74 int size=inputLength * sizeof(float);
75 cudaError_t err=cudaSuccess;
76 err=cudaMalloc((void **)&deviceInput1,size);
77 if(err!=cudaSuccess)
78 {
79     printf("GPU Memory Fatal Error = %s In File Name %s At Line No. %d.
        \nExiting ... \n",cudaGetErrorString(err),__FILE__,__LINE__);
80     cleanup();
81     exit(EXIT_FAILURE);
82 }
83
84 err=cudaMalloc((void **)&deviceInput2,size);
85 if(err!=cudaSuccess)
86 {
87     printf("GPU Memory Fatal Error = %s In File Name %s At Line No. %d.
        \nExiting ... \n",cudaGetErrorString(err),__FILE__,__LINE__);
88     cudaFree(deviceInput1);
89     cleanup();
90     exit(EXIT_FAILURE);
91 }
92
93 err=cudaMalloc((void **)&deviceOutput,size);
94 if(err!=cudaSuccess)
95 {
96     printf("GPU Memory Fatal Error = %s In File Name %s At Line No. %d.
        \nExiting ... \n",cudaGetErrorString(err),__FILE__,__LINE__);
97     cleanup();
98     exit(EXIT_FAILURE);
99 }
100
101 // copy host memory contents to device memory
102 err=cudaMemcpy(deviceInput1,hostInput1,size,cudaMemcpyHostToDevice);
103 if(err!=cudaSuccess)
104 {
105     printf("GPU Memory Fatal Error = %s In File Name %s At Line No. %d.
        \nExiting ... \n",cudaGetErrorString(err),__FILE__,__LINE__);
```

```
106     cleanup();
107     exit(EXIT_FAILURE);
108 }
109
110 err=cudaMemcpy(deviceInput2,hostInput2,size,cudaMemcpyHostToDevice);
111 if(err!=cudaSuccess)
112 {
113     printf("GPU Memory Fatal Error = %s In File Name %s At Line No. %d.
114           \nExiting ... \n",cudaGetErrorString(err),__FILE__,__LINE__);
115     cleanup();
116     exit(EXIT_FAILURE);
117 }
118
119 // cuda kernel configuration
120 dim3 DimGrid=dim3(ceil(inputLength/256.0),1,1);
121 dim3 DimBlock=dim3(256,1,1);
122 vecAdd<<<DimGrid,DimBlock>>>
123     (deviceInput1,deviceInput2,deviceOutput,inputLength);
124
125 // copy device memory to host memory
126 err=cudaMemcpy(hostOutput,deviceOutput,size,cudaMemcpyDeviceToHost);
127 if(err!=cudaSuccess)
128 {
129     printf("GPU Memory Fatal Error = %s In File Name %s At Line No. %d.
130           \nExiting ... \n",cudaGetErrorString(err),__FILE__,__LINE__);
131     cleanup();
132     exit(EXIT_FAILURE);
133 }
134
135 // results
136 int i;
137 for(i=0;i<inputLength;i++)
138 {
139     printf("%f + %f = %f\n",hostInput1[i],hostInput2[i],hostOutput[i]);
140 }
141
142 // total cleanup
143 cleanup();
144
145 return(0);
146 }
147
148 void cleanup(void)
149 {
150     // code
151
152     // free allocated device-memory
153     if(deviceInput1)
154     {
155         cudaFree(deviceInput1);
156         deviceInput1=NULL;
157     }
158
159     if(deviceInput2)
160     {
161         cudaFree(deviceInput2);
```

```
159     deviceInput2=NULL;
160 }
161
162 if(deviceOutput)
163 {
164     cudaFree(deviceOutput);
165     deviceOutput=NULL;
166 }
167
168 // free allocated host-memory
169 if(hostInput1)
170 {
171     free(hostInput1);
172     hostInput1=NULL;
173 }
174
175 if(hostInput2)
176 {
177     free(hostInput2);
178     hostInput2=NULL;
179 }
180
181 if(hostOutput)
182 {
183     free(hostOutput);
184     hostOutput=NULL;
185 }
186 }
187
```